

Jin-Soo Kim
(jinsoo.kim@snu.ac.kr)

Systems Software &
Architecture Lab.

Seoul National University

Jan. 6 – 17, 2020

Python for Data Analytics

Pandas I



Outline

- Why Pandas?
- Pandas Series
- Pandas DataFrame
- I/O in Pandas
- Time Series Data in Pandas

Why Pandas?

Limitations in NumPy

- Remember? Array slicing in NumPy

```
>>> a = np.array([[1, 2, 3], [4, 5, 6]])  
>>> a[1,:]   
array([[4, 5, 6]])  
>>> a[:,2]  
array([3, 6])  
>>> a[-1:,-2:]  
array([[5, 6]])
```

	1, 2, 3	
	4, 5, 6	

How about?

	AAPL_High	AAPL_Low
Date		
2010-01-04	214.499996	212.380001
2010-01-05	215.589994	213.249994
2010-01-06	215.230000	210.750004
2010-01-07	212.000006	209.050005
2010-01-08	212.000006	209.060005

2010-01-06 ~ 2010-01-07
사이에 발생한 data 추출?

2010년에 월별로 발생한
data를 grouping?

Limitations in NumPy (cont'd)

How about?

	AAPL_High	AAPL_Low
Date		
2010-01-04	214.499996	212.380001
2010-01-05	215.589994	213.249994
2010-01-06	215.230000	210.750004
2010-01-07	212.000006	209.050005
2010-01-08	212.000006	209.060005

	GOOG_High	GOOG_Low
Date		
2010-01-04	629.511067	624.241073
2010-01-05	627.841071	621.541045
2010-01-06	625.861078	606.361042
2010-01-07	610.001045	592.651008
2010-01-08	603.251034	589.110988

	AAPL_High	AAPL_Low	GOOG_High	GOOG_Low
Date				
2010-01-04	214.499996	212.380001	629.511067	624.241073
2010-01-05	215.589994	213.249994	627.841071	621.541045
2010-01-06	215.230000	210.750004	625.861078	606.361042
2010-01-07	212.000006	209.050005	610.001045	592.651008
2010-01-08	212.000006	209.060005	603.251034	589.110988

두 테이블의 join?

SQL and Tables (I)

- Find all instructors in Comp. Sci. dept. with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and salary > 80000;
```

Instructor relation			
ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

ID	name	dept_name	salary
83821	Brandt	Comp. Sci.	92000

SQL and Tables (2)

- For all instructors who have taught courses, find their names and the course ID of the courses they taught

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID;
```

```
select *
from instructor natural join teaches;
```

instructor

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

teaches

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010
83821	Brandt	Comp. Sci.	92000	CS-190	1	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-190	2	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-319	2	Spring	2010
98345	Kim	Elec. Eng.	80000	EE-181	1	Spring	2009

instructor ⋈ *teaches*

SQL and Tables (3)

- Group instructors in each department

```
select *  
from instructor  
group by dept_name;
```

- Find the average salary of instructors in each department

```
select dept_name, avg(salary) as avg_salary  
from instructor  
group by dept_name;
```

Instructor relation			
ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000



ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000



dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000


Data Collection for Data Analytics

- You will typically get data in one of four ways:
 1. Directly download a data file (or files) manually
 2. Query data from a database
 3. Query an API (usually web-based, these days)
 4. Scrap data from a webpage

How to perform data preprocessing in Python?

Data Format in Data Analytics





Kaggle 2017 DS & ML Survey



8,024 responses

<https://www.kaggle.com/surveys/2017>

Many Pandas Books



What is "Pandas" Module?

- panel data analysis or Python data analysis
- For building and manipulating "relational" or "tabular" data both easy and intuitive
- Built on top of NumPy (2005)
- Open source
 - Original author: Wes McKinney
 - Now part of the PyData project focused on improving Python data libraries
 - <http://pandas.pydata.org>
- `>>> import panda as pd`

Pandas History

- Developer **Wes McKinney** started working on Pandas in 2008 while at AQR Capital Management (global investment management firm)
- Need for a **high performance, flexible analysis tool for quantitative analysis on financial data**
- Before leaving AQR, he was able to convince management to allow him to open source the library
- Another AQR employee, Chang She, joined the effort in 2012 as the second major contributor to the library
- In 2015, Pandas signed on as a sponsored project of NumFOCUS, a non-profit charity in United States

Pandas Module

- Primary data structures
 - **Series** (1-dimensional)
 - **DataFrame** (2-dimensional) -- similar to *data.frame* in R
 - **Panel** (3-dimensional or more)
- Things that pandas does well
 - Easy handling of **missing data**
 - **Size mutability**: columns can be inserted and deleted (**Add & drop columns**)
 - **Powerful, flexible group by** functionality: **Groupby & aggregation**
 - **Intelligent label-based slicing, fancy indexing, and subsetting** of large data sets
 - **Intuitive merging and joining** data sets: **Join (merge) two data**
 - **Robust I/O tools** for loading data from CSV & Excel files, database, and web sources


Pandas Series

Creating Pandas Series

- 1-D array of indexed data from **Python list**

```
>>> import pandas as pd
>>> import numpy as np
>>> s = pd.Series([1,3,np.nan,6,8])
>>> s
0      1.0
1      3.0
2      NaN
3      6.0
4      8.0
dtype: float64
```

**automatic indexing
(record id/key)**



- 1-D array of indexed data from **NumPy ndarray**

```
>>> a = np.array([1,3,np.nan,6,8])
>>> s2 = pd.Series(a)
>>> s2
0      1.0
1      3.0
2      NaN
3      6.0
4      8.0
dtype: float64
```


Creating Pandas Series (cont'd)


- 1-D array of indexed data from Python dictionary

```
>>> import pandas as pd
>>> import numpy as np
>>> d = {'spam':5.99, 'egg':0.99, 'ham':3.99}
>>> s = pd.Series(d)
egg      0.99
ham      3.99
spam     5.99
dtype: float64
```

pandas.Series()

- `pd.Series([data], [index], [dtype], ...)`
 - One-dimensional ndarray with axis labels (including time series)
 - `data`: Contains data stored in Series
 - `index`: Values must be hashable and have the same length as `data` (default: `np.arange(len(data))`)
 - Non-unique index values are allowed

```
>>> a = [2, 3, 5, 8]
>>> b = ['a', 'b', 'c', 'c']
>>> s = pd.Series(a)
>>> s
0      2
1      3
2      5
3      8
dtype: int64
>>> s2 = pd.Series(a, b)
>>> s2
a      2
b      3
c      5
c      8
dtype: int64
```



Handling Missing Entries

- Series creation from dictionary
- Extracting series-index from other list

```
>>> sdata = {'Ohio':35000, \
              'Texas':71000, \
              'Oregon':16000, \
              'Utah':5000}

>>> s = pd.Series(sdata)

>>> s
Ohio      35000
Oregon    16000
Texas     71000
Utah      5000
dtype: int64
```

```
>>> states = {'California', \
              'Ohio', 'Oregon', \
              'Texas'}

>>> st = pd.Series(sdata, \
                   index=states)

>>> st
Ohio      35000.0
Texas     71000.0
Oregon    16000.0
California      NaN ← value unknown
dtype: float64
```

Checking Null Values

- `pd.isnull(obj)`
 - Return an array of Boolean indicating whether the corresponding element is missing
 - Same as `obj.isnull()`
- `pd.notnull(obj)`
 - Detect non-missing values
 - Same as `obj.notnull()`

```
>>> pd.isnull(st)
Ohio                False
Texas               False
Oregon              False
California           True
dtype: bool
```

```
>>> st.notnull()
Ohio                True
Texas               True
Oregon              True
California           False
dtype: bool
```

Pandas DataFrame


Creating Pandas DataFrame

- Dataframe is 2-D array of indexed data
 - Similar to a spreadsheet or SQL table
- Dataframe is the most commonly used pandas object
- Creating dataframe from Python dictionary

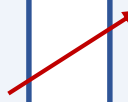
```
>>> import pandas as pd
>>> import numpy as np
>>> sales = {'Day': ['M', 'T', 'W', 'Th', 'F'], \
            'Visitors': [43, 45, 33, 43, 78], \
            'Revenue': [64, 73, 62, 64, 53]}
>>> df = pd.DataFrame(sales)
>>> df
```

**automatic indexing
(record id/key)**

column



	Day	Revenue	Visitors
0	M	64	43
1	T	73	45
2	W	62	33
3	Th	64	43
4	F	53	78



pandas.DataFrame()

- `pd.DataFrame([data], [index], [columns], [dtype], ...)`
 - The primary pandas data structure
 - Two-dimensional size-mutable, potentially heterogeneous tabular data structure with labeled axes (rows and columns)
 - `data`: ndarray, list, dictionary, or dataframe
 - `index`: index to use for resulting frame. (default: `np.arange(len(data))`)
 - `columns`: column labels to use for resulting frame

```
>>> a = {'c0':[2, 3, 5, 8], \
        'c1':[12, 76, 32, 29]}
>>> b = ['a', 'b', 'c', 'd']
>>> s = pd.DataFrame(a)
>>> s
   c0  c1
0    2  12
1    3  76
2    5  32
3    8  29
>>> s2 = pd.DataFrame(a, b)
>>> s2
   c0  c1
a    2  12
b    3  76
c    5  32
d    8  29
```

Indexing

- Row id = **key** = label = record id = **index**
- Used for
 - Accessing individual/multiple rows
 - Aligning multiple DataFrames and Series
- **`df.set_index(keys, ...)`**
 - Set the DataFrame index using existing columns
 - **Return a new DataFrame** with changed row labels (not in-place update)
 - **keys**: label or array/list of labels
e.g., `df = df.set_index(['Day', 'Revenue'])`

```
>>> df.loc[2]
Day          W
Revenue      62
Visitors     33
Name: 2, dtype: object

>>> df = df.set_index('Day')
>>> df
```

	Revenue	Visitors
Day		
M	64	43
T	73	45
W	62	33
Th	64	43
F	53	78

Accessing Rows/Columns

■ Basic operations

Operation	Syntax	Result
Select column	<code>df[col]</code>	Series
Select row by label	<code>df.loc[label]</code>	Series
Select row by integer location	<code>df.iloc[loc]</code>	Series
Slice rows	<code>df[0:2]</code>	DataFrame
Select rows by Boolean vector	<code>df[bool_vec]</code>	DataFrame

- `df.loc`: A slice object with labels [*start:stop*], both the *start* and the *stop* are **included**!

```
>>> df.loc['T']
Revenue      73
Visitors     45
Name: T, dtype: int64
>>> df.loc['M':'T'] inclusive!
      Revenue  Visitors
Day
M           64         43
T           73         45
>>> df.iloc[2:4] exclusive!
      Revenue  Visitors
Day
W           62         33
Th          64         43
```

Accessing Rows/Columns (cont'd)

```
>>> df.loc[['M', 'F']]
```

```
Revenue  Visitors
```

```
Day
```

```
M          64          43
```

```
F          53          78
```

```
>>> df['Visitors']
```

```
Day
```

```
M          43
```

```
T          45
```

```
W          33
```

```
Th         43
```

```
F          78
```

```
Name: Visitors, dtype: int64
```

```
>>> df[['Visitors', 'Revenue']]
```

```
Visitors  Revenue
```

```
Day
```

```
M          43          64
```

```
T          45          73
```

```
W          33          62
```

```
Th         43          64
```

```
F          78          53
```

```
>>> df['Visitors']['M':'W']
```

```
Day
```

```
M          43
```

```
T          45
```

```
W          33
```

```
Name: Visitors, dtype: int64
```

Boolean Indexing

```
>>> df[[True,False,False,True,False]]
```

	Revenue	Visitors
--	---------	----------

Day		
-----	--	--

M	64	43
---	----	----

Th	64	43
----	----	----

```
>>> df[df['Revenue'] > 65]
```

	Revenue	Visitors
--	---------	----------

Day		
-----	--	--

T	73	45
---	----	----

```
>>> df[(df['Revenue'] > 50) & \
```

```
      (df['Visitors'] > 50)]
```

	Revenue	Visitors
--	---------	----------

Day		
-----	--	--

F	53	78
---	----	----

```
>>> df[(df['Revenue'] > 50) | \
```

```
      (df['Visitors'] > 50)]
```

	Revenue	Visitors
--	---------	----------

Day		
-----	--	--

T	73	45
---	----	----

F	53	78
---	----	----

```
>>> df[df > 50]
```

	Revenue	Visitors
--	---------	----------

Day		
-----	--	--

M	64	NaN
---	----	-----

T	73	NaN
---	----	-----

W	62	NaN
---	----	-----

Th	64	NaN
----	----	-----

F	53	78.0
---	----	------

Column Manipulation

- Change the order of columns

```
>>> df2 = pd.DataFrame(df, columns=['Visitors','Revenue'])
>>> df2
```

	Visitors	Revenue
Day		
M	43	64
T	45	73
W	33	62
Th	43	64
F	78	53

- Add a new column:
 - NaN are filled to added column values

```
>>> df3 = pd.DataFrame(df, columns=['Visitors','Revenue','Debt'])
>>> df3
```

	Visitors	Revenue	Debt
Day			
M	43	64	NaN
T	45	73	NaN
W	33	62	NaN
Th	43	64	NaN
F	78	53	NaN

Column Manipulation (cont'd)

- Delete an existing column

- Using `del` (delete in-place)

```
>>> del df3['Debt']
>>> df3
```

	Visitors	Revenue
Day		
M	43	64
T	45	73
W	33	62
Th	43	64
F	78	53

axis=0

axis=1

Day	Visitors	Revenue	Debt
M	43	64	NaN
T	45	73	NaN
W	33	62	NaN
Th	43	64	NaN
F	78	53	NaN

- Using `df.drop()` (return a new df)

```
>>> df3.drop('Debt', axis=1)
```

	Visitors	Revenue
Day		
M	43	64
T	45	73
W	33	62
Th	43	64
F	78	53

Row Manipulation

- Add a new row

```
>>> df2.loc['S'] = [92, 87]
>>> df2
```

Visitors	Revenue
Day	
M	43 64
T	45 73
W	33 62
Th	43 64
F	78 53
S	92 87

axis=0

axis=1

Day	Visitors	Revenue
M	43	64
T	45	73
W	33	62
Th	43	64
F	78	53
S	92	87

- Delete an existing row

- Using `df.drop()`
(return a new df)

```
>>> df2.drop('S', axis=0)
```

Visitors	Revenue	Debt
Day		
M	43	64
T	45	73
W	33	62
Th	43	64
F	78	53

Rename Row/Column

- `df.rename([index], [columns], [inplace], ...)`
 - Rename any index, row or column
 - A part of rows or columns can be altered
 - `index`: dict. for changing row indexes
 - `columns`: dict. for changing column indexes
 - `inplace`: If True, `df` is updated in place. Otherwise, return a new `df` (default: False)

```
>>> newr = {'M':'Mo', 'T':'Tu'}
>>> df.rename(index=newr)
```

	Revenue	Visitors
Day		
Mo	64	43
Tu	73	45
W	62	33
Th	64	43
F	53	78

```
>>> newc = {'Revenue':'Rev.'}
>>> df.rename(columns=newc)
```

	Rev.	Visitors
Day		
M	64	43
T	73	45
W	62	33
Th	64	43
F	53	78

Common Statistical Functions

Method	Description
<code>count()</code>	Number of non-null observations
<code>sum()</code>	Sum of values
<code>mean()</code>	Mean of values
<code>median()</code>	Arithmetic median of values
<code>min()</code>	Minimum
<code>max()</code>	Maximum
<code>std()</code>	Bessel-corrected sample standard deviation
<code>var()</code>	Unbiased variance
<code>skew()</code>	Sample skewness (3rd moment)
<code>kurt()</code>	Sample kurtosis (4th moment)
<code>quantile()</code>	Sample quantile (value at %)
<code>apply()</code>	Generic apply
<code>cov()</code>	Unbiased covariance
<code>corr()</code>	Correlation

- Applicable both **Series** and **DataFrame** objects

```
>>> df[df['Revenue']>60].count()
Revenue      4
Visitors     4
dtype: int64
>>> df['Revenue'].mean()
63.2
>>> df.cov()
           Revenue  Visitors
Revenue    50.70    -81.35
Visitors   -81.35    295.80
>>> df.corr()
           Revenue  Visitors
Revenue    1.000000 -0.664285
Visitors  -0.664285  1.000000
```


Iteration over Rows

- `df.iterrows()`
 - Iterate over rows of DataFrame as (index, Series) pairs

```
>>> df = pd.DataFrame(data=[[1,2,3],[4,5,6],[7,8,9]], columns=['A','B','C'])
>>> df
```

	A	B	C
0	1	2	3
1	4	5	6
2	7	8	9

```
>>> for index, row in df.iterrows():
...     print(row['A'], row['B'])
1 2
4 5
7 8
```

```
>>> for index, row in df.iterrows():
...     print(row)
```

```
A    1
B    2
C    3
Name: 0, dtype: int64
A    4
B    5
C    6
Name: 1, dtype: int64
A    7
B    8
C    9
Name: 2, dtype: int64
```

} **Series**

Groupby and Aggregation (I)

■ Example DataFrame object

- 2 categorical values: A, B
- 2 numerical values: C, D

```
>>> df = pd.DataFrame({'A': ['foo', 'bar', 'foo', 'bar', 'foo', 'bar', 'foo', 'foo'], \
                        'B': ['one', 'one', 'two', 'three', 'two', 'two', 'one', 'three'], \
                        'C': np.random.randn(8), \
                        'D': np.random.randn(8)})
```

```
>>> df
```

	A	B	C	D
0	foo	one	-0.578235	0.193109
1	bar	one	1.312911	0.576292
2	foo	two	0.628944	0.484595
3	bar	three	0.206827	0.810682
4	foo	two	1.584507	1.153200
5	bar	two	-0.367555	-0.703818
6	foo	one	-0.017915	-1.297967
7	foo	three	0.337489	1.565021

```
>>> df[df['A'] == 'foo']
```

	A	B	C	D
0	foo	one	-0.578235	0.193109
2	foo	two	0.628944	0.484595
4	foo	two	1.584507	1.153200
6	foo	one	-0.017915	-1.297967
7	foo	three	0.337489	1.565021

```
>>> df[df['A'] == 'bar']
```

	A	B	C	D
1	bar	one	1.312911	0.576292
3	bar	three	0.206827	0.810682
5	bar	two	-0.367555	-0.703818

Groupby and Aggregation (2)

- `df.groupby([by], [axis], ...)`
 - Used to group large amounts of data and compute operations on these groups
 - `by`: label, function, a list of labels, ...
(Used to determine the groups)
 - `axis`: 0 or 'index' for rows, 1 or 'columns' for columns (default: 0)
- Aggregation stat functions after grouping
 - `mean()`, `sum()`, `median()`, `var()`, etc.

```
>>> g = df.groupby('A')
>>> g.mean()
```

	C	D
A		
bar	0.384061	0.227719
foo	0.390958	0.419592

```
>>> g.sum()
```

	C	D
A		
bar	1.152183	0.683156
foo	1.954789	2.097958

```
>>> g.corr()
```

		C	D
A			
bar	C	1.000000	0.661077
	D	0.661077	1.000000
foo	C	1.000000	0.493865
	D	0.493865	1.000000

Groupby and Aggregation (3)

- Get a group's contents
- Printing the groups

```
>>> g.get_group('bar')
   B          C          D
1  one  1.312911  0.576292
3 three  0.206827  0.810682
5  two -0.367555 -0.703818
```

```
>>> for key, item in g:
...     print(key)
...     print(g.get_group(key))
bar
   B          C          D
1  one  1.312911  0.576292
3 three  0.206827  0.810682
5  two -0.367555 -0.703818
foo
   B          C          D
0  one -0.578235  0.193109
2  two  0.628944  0.484595
4  two  1.584507  1.153200
6  one -0.017915 -1.297967
7 three  0.337489  1.565021
```

Groupby and Aggregation (4)

- Describing a group

```
>>> g.describe()
C
count      mean      std      min      25%      50%      75%
A
bar      3.0      0.384061      0.854137      -0.367555      -0.080364      0.206827      0.759869
foo      5.0      0.390958      0.804762      -0.578235      -0.017915      0.337489      0.628944

D
max count      mean      std      min      25%      50%
A
bar      1.312911      3.0      0.227719      0.815202      -0.703818      -0.063763      0.576292
foo      1.584507      5.0      0.419592      1.101784      -1.297967      0.193109      0.484595

75%      max
A
bar      0.693487      0.810682
foo      1.153200      1.565021
```

- Grouping by multiple columns

```
>>> gm = df.groupby(['A', 'B'])
>>> gm.mean()
A  B      C      D
bar one      1.312911      0.576292
    three      0.206827      0.810682
    two      -0.367555      -0.703818
foo one      -0.298075      -0.552429
    three      0.337489      1.565021
    two      1.106725      0.818898
```

Merging (Joining)

- `pd.merge(left, right, [how], [on], [left_on], [right_on], [left_index], [right_index], ...)`
 - Merge DataFrame objects with database-style join
 - `left`: DataFrame
 - `right`: Object to merge with
 - `how`: join type -- 'left', 'right', 'outer', or 'inner' (default: 'inner')
 - `on`: column to join on (label or list) -- must be found on both DataFrames
 - `left_on` (or `right_on`): column to join on in the left (or right) DataFrame
 - `left_index` (or `right_index`): if True, use the index from the left (or right) DataFrame

```
result =  
    pd.merge(left, right, on='key')
```

left				right				Result					
	A	B	key		C	D	key		A	B	key	C	D
0	A0	B0	K0	0	C0	D0	K0	0	A0	B0	K0	C0	D0
1	A1	B1	K1	1	C1	D1	K1	1	A1	B1	K1	C1	D1
2	A2	B2	K2	2	C2	D2	K2	2	A2	B2	K2	C2	D2
3	A3	B3	K3	3	C3	D3	K3	3	A3	B3	K3	C3	D3

Merging (Joining) (cont'd)

- Merge two DataFrames by their own index


```
dfxy = pd.merge(dfx, dfy, left_index=True, right_index=True)
```

dfx			dfy			dfxy				
	AAPL_High	AAPL_Low		GOOG_High	GOOG_Low		AAPL_High	AAPL_Low	GOOG_High	GOOG_Low
Date			Date			Date				
2010-01-04	214.499996	212.380001	2010-01-04	629.511067	624.241073	2010-01-04	214.499996	212.380001	629.511067	624.241073
2010-01-05	215.589994	213.249994	2010-01-05	627.841071	621.541045	2010-01-05	215.589994	213.249994	627.841071	621.541045
2010-01-06	215.230000	210.750004	2010-01-06	625.861078	606.361042	2010-01-06	215.230000	210.750004	625.861078	606.361042
2010-01-07	212.000006	209.050005	2010-01-07	610.001045	592.651008	2010-01-07	212.000006	209.050005	610.001045	592.651008
2010-01-08	212.000006	209.060005	2010-01-08	603.251034	589.110988	2010-01-08	212.000006	209.060005	603.251034	589.110988

Merging (Joining) Types


- Inner join ('inner')
 - Return only the rows in which the left table have matching keys in the right table
- Outer join ('outer')
 - Returns all rows from both tables, join records from the left which have matching keys in the right table.
- Left outer join ('left')
 - Return all rows from the left table, and any rows with matching keys from the right table.
- Right outer join ('right')
 - Return all rows from the right table, and any rows with matching keys from the left table.

how='inner'




natural join

how='outer'




full outer join

how='left'



left outer join

how='right'



right outer join

Merging (Joining) Example

df1

	id	name
0	1	Alice
1	2	Bob
2	3	Charlie
3	4	David
4	5	Emily

df2

	id	country
0	2	Korea
1	4	US
2	5	UK
3	6	Italy

inner

	id	name	country
0	2	Bob	Korea
1	4	David	US
2	5	Emily	UK

outer

	id	name	country
0	1	Alice	NaN
1	2	Bob	Korea
2	3	Charlie	NaN
3	4	David	US
4	5	Emily	UK
5	6	NaN	Italy

`pd.merge(df1, df2)`

left

	id	name	country
0	1	Alice	NaN
1	2	Bob	Korea
2	3	Charlie	NaN
3	4	David	US
4	5	Emily	UK

right

	id	name	country
0	2	Bob	Korea
1	4	David	US
2	5	Emily	UK
3	6	NaN	Italy

I/O in Pandas

I/Os for Pandas DataFrame

- A collection of convenient I/O functions supporting various file formats

<code>to_csv()</code>	<code>to_excel()</code>	<code>to_hdf()</code>	<code>to_sql()</code>	<code>to_json()</code>	<code>to_html()</code>
<code>read_csv()</code>	<code>read_excel()</code>	<code>read_hdf()</code>	<code>read_sql()</code>	<code>read_json()</code>	<code>read_html()</code>

- (cf.) HDF (Hierarchical Data Format): Standardized file format for scientific data
- From Pandas DataFrame to CSV file: `pd.to_csv(path)`
- From CSV file to Pandas DataFrame: `pd.read_csv(path)`

Reading a CSV File

```
>>> df = pd.read_csv('GOOGL.csv')
>>> df.iloc[0:9]
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	2009-05-22	198.528534	199.524521	196.196198	196.946945	196.946945	3433700
1	2009-05-26	196.171173	202.702698	195.195190	202.382385	202.382385	6202700
2	2009-05-27	203.023026	206.136139	202.607605	202.982986	202.982986	6062500
3	2009-05-28	204.544540	206.016022	202.507507	205.405411	205.405411	5332200
4	2009-05-29	206.261261	208.823822	205.555557	208.823822	208.823822	5291100
5	2009-06-01	209.574570	215.015015	209.474472	213.493500	213.493500	6638100
6	2009-06-02	213.338333	215.195190	211.911911	214.414413	214.414413	5241900
7	2009-06-03	213.213211	216.446442	212.212219	216.041046	216.041046	7058500
8	2009-06-04	217.867874	220.840836	217.467468	220.360367	220.360367	7268900
9	2009-06-05	222.757751	223.893890	219.949951	222.382385	222.382385	7354200

GOOGL.csv - Excel							
파일	홈	삽입	페이지 레이아웃	수식	데이터	검토	보기
K8							
	A	B	C	D	E	F	G
1	Date	Open	High	Low	Close	Adj Close	Volume
2	2009-05-22	198.52853	199.5245	196.1962	196.9469	196.9469	3433700
3	2009-05-26	196.17117	202.7027	195.1952	202.3824	202.3824	6202700
4	2009-05-27	203.02303	206.1361	202.6076	202.983	202.983	6062500
5	2009-05-28	204.54454	206.016	202.5075	205.4054	205.4054	5332200
6	2009-05-29	206.26126	208.8238	205.5556	208.8238	208.8238	5291100
7	2009-06-01	209.57457	215.015	209.4745	213.4935	213.4935	6638100
8	2009-06-02	213.33833	215.1952	211.9119	214.4144	214.4144	5241900
9	2009-06-03	213.21321	216.4464	212.2122	216.041	216.041	7058500
10	2009-06-04	217.86787	220.8408	217.4675	220.3604	220.3604	7268900
11	2009-06-05	222.75775	223.8939	219.95	222.3824	222.3824	7354200
12	2009-06-08	219.96997	220.6807	217.2773	219.6046	219.6046	6191200

pandas.read_csv()

- `pd.read_csv(filepath, [sep], [header], [names], [index_col], [encoding], ...)`
 - Read a comma-separated values (csv) file
 - *filepath*: any valid string path. The string could be a URL.
 - *sep* (or *delimiter*): delimiter to use (default: ',')
 - *header*: row number(s) to use as the column names
 - *names*: list of column names to use
 - *index_col*: column(s) to use as the row labels of the Data Frame
 - *encoding*: encoding to use (default: 'utf-8')

```
>>> df2 = pd.read_csv('GOOGL.csv', names=['A','B','C','D','E','F'])  
>>> df2 = pd.read_csv('mydf.csv', sep=':')
```

DataFrame.to_csv()

- `df.to_csv(filepath, [sep], [columns], [header], [index], [encoding], ...)`
 - Write DataFrame to a comma-separated values (csv) file
 - `filepath`: any valid string path.
 - `sep` (or `delimiter`): delimiter to use (default: ',')
 - `columns`: columns to write
 - `header`: write out the column names (default: True)
 - `index`: write row names (default: True)
 - `encoding`: encoding to use (default: 'utf-8')

```
>>> df.to_csv('mydf.csv', sep='\t')
>>> df.to_csv('dataset.csv', sep='\t', encoding='utf-8')
```