

Jin-Soo Kim
(jinsoo.kim@snu.ac.kr)

Systems Software &
Architecture Lab.

Seoul National University

Jan. 6 – 17, 2020

Python for Data Analytics

Pandas II



Outline

- Why Pandas?
- Pandas Series
- Pandas DataFrame
- I/O in Pandas
- Time Series Data in Pandas

Times Series Data in Pandas

Time Series Data

- How to analyze time series data?

- Sample time series data

2011-01-01 00:00:00	-0.131254
2011-01-01 01:00:00	0.068876
2011-01-01 02:00:00	-0.207636
2011-01-01 03:00:00	1.388030
Timestamp Index → 2011-01-01 04:00:00	0.937158

- Time series data is the data with the timestamp index

- How to parse time series information from various sources and formats?
- How to generate sequences of fixed-frequency dates and time spans
- How to manipulate and convert date times with timezone information?
- How to group data by time?
- ...

Python datetime Module

- `datetime.datetime` class: a combination of date and time
 - year, month, day, hour, minute, second, microsecond, tzinfo
- `datetime.now()`: return the current local datetime

```
>>> import datetime as dt
>>> now = dt.datetime.now()
>>> print(now)
2020-01-04 01:01:35.066589
>>> newyear = dt.datetime(2020, 1, 1)
>>> print(newyear)
2020-01-01 00:00:00
>>> print(now - newyear)
3 days, 1:01:35.066589
```

NumPy datetime64 Type

- NumPy supports datetime functionality with the data type called **'datetime64'**
 - No timezone support

```
>>> import numpy as np
>>> now = np.datetime64('now')
>>> print(now)
2020-01-03T16:39:44
>>> np.arange('2020-01', '2020-07', dtype='datetime64[M]')
array(['2020-01', '2020-02', '2020-03', '2020-04', '2020-05', '2020-06'],
      dtype='datetime64[M]')
>>> newyear = np.datetime64('2020-1-1')
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: Error parsing datetime string "2020-1-1" at position 5
```

Converting to Datetime

- Pandas supports extensive capabilities and features for working with time series data based on NumPy datetime64.
- `pd.to_datetime(arg, ...)`
 - Convert argument to datetime.
 - Return type can be a DatetimeIndex, Series, or Timestamp
 - *arg*: integer, float, string, datetime, list, tuple, I-D array, Series

```
>>> t = np.array(['1/8/2020', '2020/1/9', '20200110', '2020-1-13', '2020 1 14',  
                  '2020, 1, 15', 'Jan. 16 2020', '17 Jan 2010'])  
>>> pd.to_datetime(t)  
DatetimeIndex(['2020-01-08', '2020-01-09', '2020-01-10', '2020-01-13',  
               '2020-01-14', '2020-01-15', '2020-01-16', '2010-01-17'],  
              dtype='datetime64[ns]', freq=None)
```

Generating DatetimeIndex

- `pd.date_range([start], [end], [periods], [freq], ...)`
 - Return a fixed frequency datetime index
 - *start*: Left bound for generating dates
 - *end*: Right bound for generating dates
 - *periods*: the number of datetime to generate
 - *freq*: the time interval between consecutive datetime values (default: 'D')

Freq string	Description	Freq string	Description
'D'	One absolute day	'M'	Calendar month end
'H'	One hour	'MS'	Calendar month begin
'T' or 'min'	One minute	'BM'	Business month end
'S'	One second	'BMS'	Business month begin
'B'	Business day (weekday)	'WOM-2THU'	Second Thursday of the month
'W'	One week	'1h30min'	One and half hour

date_range() Examples (I)

- Default: everyday

```
>>> pd.date_range('2020-1-6', '2020-1-17')
DatetimeIndex(['2020-01-06', '2020-01-07', '2020-01-08', '2020-01-09',
               '2020-01-10', '2020-01-11', '2020-01-12', '2020-01-13',
               '2020-01-14', '2020-01-15', '2020-01-16', '2020-01-17'],
              dtype='datetime64[ns]', freq='D')
```

- 14 days since 2020-1-6

```
>>> pd.date_range('2020 1 6', periods=14)
DatetimeIndex(['2020-01-06', '2020-01-07', '2020-01-08', '2020-01-09',
               '2020-01-10', '2020-01-11', '2020-01-12', '2020-01-13',
               '2020-01-14', '2020-01-15', '2020-01-16', '2020-01-17',
               '2020-01-18', '2020-01-19'],
              dtype='datetime64[ns]', freq='D')
```

date_range() Examples (2)

- Just weekdays

```
>>> pd.date_range('2020 1 1', '2020/1/20', freq='B')
DatetimeIndex(['2020-01-01', '2020-01-02', '2020-01-03', '2020-01-06',
              '2020-01-07', '2020-01-08', '2020-01-09', '2020-01-10',
              '2020-01-13', '2020-01-14', '2020-01-15', '2020-01-16',
              '2020-01-17', '2020-01-20'],
              dtype='datetime64[ns]', freq='B')
```

- Every Sunday

```
>>> pd.date_range('2020-1-1', '2020-3-1', freq='W-SUN')
DatetimeIndex(['2020-01-05', '2020-01-12', '2020-01-19', '2020-01-26',
              '2020-02-02', '2020-02-09', '2020-02-16', '2020-02-23',
              '2020-03-01'],
              dtype='datetime64[ns]', freq='W-SUN')
```

date_range() Examples (3)

- First business day every two months

```
>>> pd.date_range('2020-1-1', '2020-12-31', freq='2BMS')
DatetimeIndex(['2020-01-01', '2020-03-02', '2020-05-01', '2020-07-01',
               '2020-09-01', '2020-11-02'],
              dtype='datetime64[ns]', freq='2BMS')
```

- Every one and half hour

```
>>> pd.date_range('2020-1-6 9:30', periods=7, freq='1h30min')
DatetimeIndex(['2020-01-06 09:30:00', '2020-01-06 11:00:00',
               '2020-01-06 12:30:00', '2020-01-06 14:00:00',
               '2020-01-06 15:30:00', '2020-01-06 17:00:00',
               '2020-01-06 18:30:00'],
              dtype='datetime64[ns]', freq='90T')
```

Finding the Day of the Week

- `pd.DatetimeIndex.day_name(self, ...)`
 - Return the day names of the `DatetimeIndex`

```
>>> idx = pd.date_range(start='2020-01-01', freq='D', periods=3)
>>> idx
DatetimeIndex(['2020-01-01', '2020-01-02', '2020-01-03'], dtype='datetime64[ns]', freq='D')

>>> idx.day_name()
Index(['Wednesday', 'Thursday', 'Friday'], dtype='object')

>>> war = pd.date_range(start='1950-6-25', freq='D', periods=3)
>>> war.day_name()
Index(['Sunday', 'Monday', 'Tuesday'], dtype='object')
```

Creating Time Series Data

- Create a range of DatetimeIndex object

```
>>> ts = pd.date_range('1/1/2020', periods=4, freq='2H')
>>> ts
DatetimeIndex(['2020-01-01 00:00:00', '2020-01-01 02:00:00',
               '2020-01-01 04:00:00', '2020-01-01 06:00:00'],
              dtype='datetime64[ns]', freq='2H')
```

- Use the DatetimeIndex object for Pandas Series or DataFrame index

```
>>> s=pd.Series(np.random.randn(len(ts)), index=ts)
>>> s
2020-01-01 00:00:00    0.872481
2020-01-01 02:00:00   -0.258136
2020-01-01 04:00:00    1.151232
2020-01-01 06:00:00    1.787104
Freq: 2H, dtype: float64
```

Example: Pandas Time Series Data (I)

- Create a dataframe: input dataset = *<timestamp, access count>*

```
import datetime as dt
import pandas as pd
data = {'date': ['2020-01-01 08:47:05.069722',
                '2020-01-01 18:47:05.119994',
                '2020-01-02 08:47:05.178768',
                '2020-01-02 13:47:05.230071',
                '2020-01-02 18:47:05.230071',
                '2020-01-02 23:47:05.280592',
                '2020-01-03 08:47:05.332662',
                '2020-01-03 18:47:05.385109',
                '2020-01-04 08:47:05.436523',
                '2020-01-04 18:47:05.486877'],
        'counts': [34, 25, 26, 15, 15, 14, 26, 25, 62, 41]}
df = pd.DataFrame(data)
df
```

	date	counts
0	2020-01-01 08:47:05.069722	34
1	2020-01-01 18:47:05.119994	25
2	2020-01-02 08:47:05.178768	26
3	2020-01-02 13:47:05.230071	15
4	2020-01-02 18:47:05.230071	15
5	2020-01-02 23:47:05.280592	14
6	2020-01-03 08:47:05.332662	26
7	2020-01-03 18:47:05.385109	25
8	2020-01-04 08:47:05.436523	62
9	2020-01-04 18:47:05.486877	41

Example: Pandas Time Series Data (2)

- Convert `df['date']` from string to datetime

```
df['date'] = pd.to_datetime(df.date)
```

- Set `df['date']` as the index

```
df = df.set_index('date')  
df
```

counts	
date	
2020-01-01 08:47:05.069722	34
2020-01-01 18:47:05.119994	25
2020-01-02 08:47:05.178768	26
2020-01-02 13:47:05.230071	15
2020-01-02 18:47:05.230071	15
2020-01-02 23:47:05.280592	14
2020-01-03 08:47:05.332662	26
2020-01-03 18:47:05.385109	25
2020-01-04 08:47:05.436523	62
2020-01-04 18:47:05.486877	41

Example: Pandas Time Series Data (3)

- View data in 2020

```
df['2020']
```

counts	
date	
2020-01-01 08:47:05.069722	34
2020-01-01 18:47:05.119994	25
2020-01-02 08:47:05.178768	26
2020-01-02 13:47:05.230071	15
2020-01-02 18:47:05.230071	15
2020-01-02 23:47:05.280592	14
2020-01-03 08:47:05.332662	26
2020-01-03 18:47:05.385109	25
2020-01-04 08:47:05.436523	62
2020-01-04 18:47:05.486877	41

- View data in January 2020

```
df['2020-01']
```

counts	
date	
2020-01-01 08:47:05.069722	34
2020-01-01 18:47:05.119994	25
2020-01-02 08:47:05.178768	26
2020-01-02 13:47:05.230071	15
2020-01-02 18:47:05.230071	15
2020-01-02 23:47:05.280592	14
2020-01-03 08:47:05.332662	26
2020-01-03 18:47:05.385109	25
2020-01-04 08:47:05.436523	62
2020-01-04 18:47:05.486877	41

Example: Pandas Time Series Data (4)

- Observations after Jan. 3, 2020
- Observations between Jan. 1 - 2

```
df[df.datetime(2020, 1, 3):]
```

counts	
date	
2020-01-03 08:47:05.332662	26
2020-01-03 18:47:05.385109	25
2020-01-04 08:47:05.436523	62
2020-01-04 18:47:05.486877	41

```
df['1/1/2020':'1/2/2020']
```

counts	
date	
2020-01-01 08:47:05.069722	34
2020-01-01 18:47:05.119994	25
2020-01-02 08:47:05.178768	26
2020-01-02 13:47:05.230071	15
2020-01-02 18:47:05.230071	15
2020-01-02 23:47:05.280592	14

Example: Pandas Time Series Data (5)

- Mean value of counts per day

```
df.resample('D').mean()
```

counts	
date	
2020-01-01	29.5
2020-01-02	17.5
2020-01-03	25.5
2020-01-04	51.5

- Total value of counts per day

```
df.resample('D').sum()
```

counts	
date	
2020-01-01	59
2020-01-02	70
2020-01-03	51
2020-01-04	103

Example: Pandas Time Series Data (6)

- Truncate observations after Jan. 3, 2020

```
df.truncate(after='1/3/2020')
```

counts	
date	
2020-01-01 08:47:05.069722	34
2020-01-01 18:47:05.119994	25
2020-01-02 08:47:05.178768	26
2020-01-02 13:47:05.230071	15
2020-01-02 18:47:05.230071	15
2020-01-02 23:47:05.280592	14

- Plot the total counts per day

```
import matplotlib.pyplot as plt  
df.resample('D').sum().plot()  
plt.show()
```

